


Introduction to Data Science Course Project

Cars – price prediction



WordCloud

Data Collection

My source dataset has been acquired from the website Kaggle.com. This dataset has been created by the user name “Austin Reese” who scraped Craigslist.com for several months. This dataset is composed of information about used cars that have been posted on craigslist.com by users around the US. Craigslist.com is a wonderful place to collect this kind of data since my project is dealing with price prediction and to observe the correction of the price I decided to have a dataset that it is not from an official/ business is related to this field.

Data Format Description

The dataset contains over 500K lines of items and has more than 18 categories. This set of data is being activated every few months (according to its owner Austin Reese) and recorded the data from craigslist. The entire dataset is saved as a CSV file and keeps growing every time. My topic is to observe price prediction and check the accuracy of used cars that have been advertised on craigslist. The amount of data will make it as accurate as possible by using some ML models. I have removed most of the attributes since they have nothing to do with price prediction such as 'url', 'image' and 'region'. The size of the dataset is big (more than 1.5GB) so I decided to use only a third of it, later on, I should add some more information if the results will be off than my expectation.

Cleaning the dataset and remove nan information have been done deeply with some general insights. I have noticed that there are many nan rows more particularly in attributes that have a direct impact on the price for instance 'year' and 'condition'. All nan information has been changed by using two methods:

1. The condition using odometer information to complete the status condition of the car. This method allowed me to guess the condition status of each nan fill in the 'condition' attribute by using the mean function. For example, the mean result of excellent condition was 108300 (rounded) so each car that is equal to or has more than that will be named as 'excellent' and so on.

2. Prices that were more than the average price for used cars have been removed, as well as year model that was too old and very hard to predict since those cars belong to a different category of sale (antique).

3. And finally, all nan values have been filled with the 'ffill' method or dropna() function.

Finally, I have found the most relevant attributes to my project that are: Finally, I have found the most relevant attributes to my project that are:

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 113341 entries, 0 to 113340
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 113341 non-null  int64
1   year                 113127 non-null  float64
2   manufacturer         108349 non-null  object
3   model                111735 non-null  object
4   condition            63930 non-null   object
5   cylinders            68844 non-null   object
6   fuel                 112591 non-null   object
7   odometer             92688 non-null   float64
8   transmission         112497 non-null   object
9   drive                81049 non-null   object
10  type                 83113 non-null   object
11  paint_color          77291 non-null   object
12  state                113341 non-null   object
dtypes: float64(2), int64(1), object(10)
memory usage: 12.1+ MB
```


Descriptive Statistics

As I mentioned at first my project is dealing with prediction and in particular price prediction of used cars. The dataset from Craigslist has tons of information about used cars some are out of range (can be easily observed by the price – condition – year mismatch) and those were removed to work on a dataset that is accurate as much as possible. Working

on this dataset and this kind of project the main factor is the price and we want to check that if the price that was given by the owners in this case in Craigslist post is accurate and matches the prediction. This can only be done by learning from many other examples.

To start to analyze my dataset I checked the range and the distribution of some of my main attributes


Posts per state in the US:



Popular car over the dataset:


With the highest volume of posts “Ford” is the most popular. This fact should give us the factor the ‘Ford’ value is important in this observation and should be taken into consideration while working on prediction.

In this general plot we can see that the state of TX has the highest amount of posts over Craigslist, next I will use this information to learn if this case has any connection between price and




Price mean:

From this observation we can learn that the range of used cars prices is between 10K to 30K. from this plot we can see that 'price' attribute has a wide distribution. There is a peak in the price range between 10k -18K which means that most of the cars being sold on Craig list are in good




Odometer mean evaluation:

Another factor/ category that is very important when we work with used cars is the fact of how good is the car in other words how much use has been done with a particular car. From this distribution we can learn that used cars tend to be driven somewhere between 50K to 150K depends on the type, and year.



Measure car condition:

When users advertise their cars on Craigslist they want to say something about the condition of the car, this status can contribute in a different on the sale such as increase more buyers to show some interest. Also, as mentioned above this column was lack of information in many rows so the best way to complete it was according to the mileage that a cars has been driven. In this graph we can see that excellent condition is the most popular status (makes sense when we talk about used cars).



Data Analysis, Visualization, and Insights

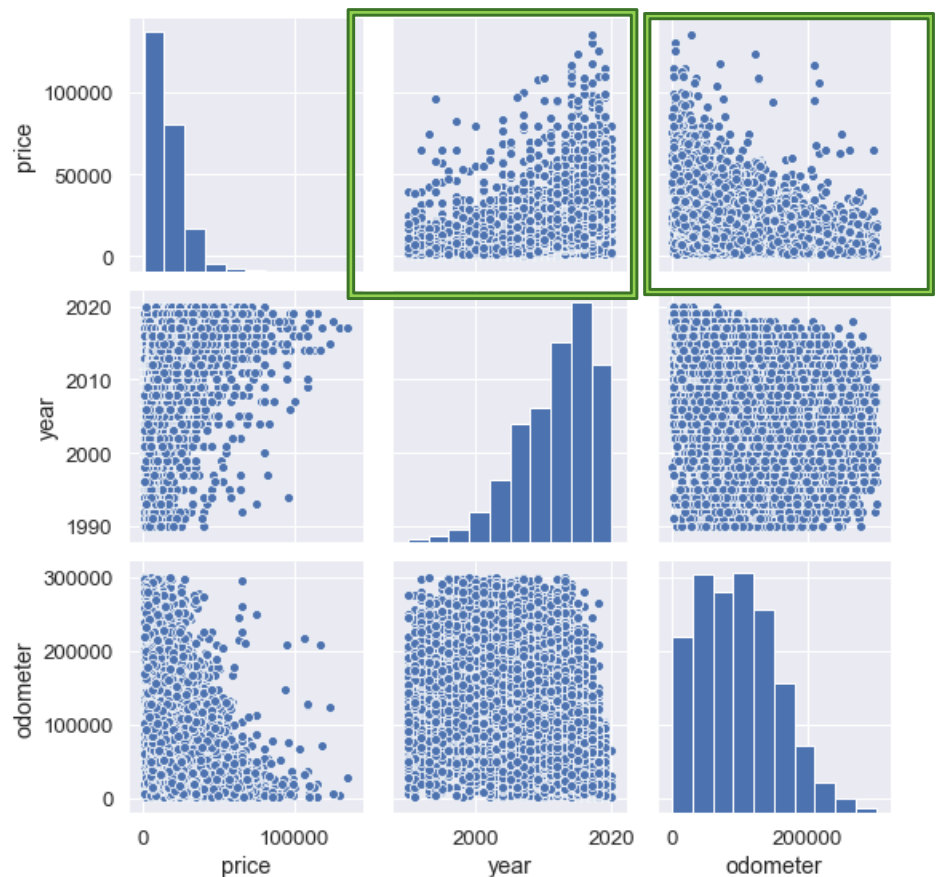
Many interesting insights can be derived from this data, such as how many cars people hold in each state? What is the most popular car and models that is been sold over Craigslist? Is Craigslist a better platform than other digital media? How to predict if the determent price is accurate. What type of engine people like? Is it different among other states? In my

project I will try to answer most of the questions above but, the main idea behind this is to predict current prices. The question is “does the price match the condition of the car?” or “What impact on the price?” When a seller tries to sell his car, he makes a guess and determines his price according to the same car that is already advertised over the media. I will try to predict how accurate is that determination about what people have on craigslist.

To understand how the prices are determined I need to see what affects the price. For this, I will explore some attributes concerning the price. Also, from my own experience I can tell that condition, odometer and year have a great impact on the price and I should start my data exploration there.


Explore features that could impact the price:

Using seaborn plot in order to explore the relation between price, odometer and year. From this graph we can easily say that the higher level of odometer the less price as well as the older car the lower price. So in fact price is related to how a specific car has been used and what is the year.




Odometer Vs Condition:

Based on the boxplot graph we can see that more than 50% of the 'good' condition cars have done more than 100K miles, however, people that tend to sell new cars drive much less where new cars consider as no older than 2 years. Most of the 'fair' condition which is ranked lowest in the condition ranking have driven more than 200K these cars are somewhere between 1990 to 2000.




Price Vs Condition:


According to this box plot, we can tell that (as expected) the better condition the higher the price, however, there is slightly different in the 'new' condition where some of the new cars are low than the total advertised cars this could be because of the model or any other variable that will be explored next. In general, this graph determines condition and price where the average is somewhere between 40 – 50% which is the number as expected. From here we cannot make a strict determination so we need to explore some more factors how prices have determined by its sellers.



With the scatter plot below we can see that most of the cars in the range of \$ 0 – 20, 0000 have in average 50,000 miles. This graph allows me to put more emphasize on Odometer Vs Price when prediction



As well as the graph bellow- using heatmap to observe the correlation between the three factors (price, year and odometer) these are my main factors that will be taking into calculation.



Future Plans

1. My next tasks are to start calculating the price prediction using some Regression models such as Linear Regression and XGBoost. These two models will help me to get the results of how the price is predicted in this dataset vs the actual price. With some error corrections, I can get the closest results.
2. Also, I will try to add some more rows into my calculation (used only 1/3 of the overall data) to see if there are some changes and differences when we add some more data.
3. Explore more models and check if they can fit into this experiment.
4. My next goal after predicting the price I will try to decrease the error factor as much as possible by trying more prediction models.
5. lastly I will check more parameters/ variables in order to check their relation to my price prediction assumption these factors will be: Drive, Fuel and so on.

Reference:

1. Austin Reese (2020, March). Used Cars dataset: Vehicles listing from Craigslist. version 5. 14, April 2020.
<https://www.kaggle.com/austinreese/craigslist-carstrucks-data>
2. XGBoost:
https://xgboost.readthedocs.io/en/latest/python/python_api.html
3. Linear Regression:
https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html